

DECODING ROAD SAFETY: A MULTIDIMENSIONAL ANALYTIC STUDY

Identifying Risk Factors, Hotspots, and Temporal Patterns for Proactive Intervention

SECTOR

Transport & Public Safety / Urban Analytics

Team Members-

- Ganga Raghuwanshi
- Krishna Devan
- Siddhant Raikar
- Arjun Singh
- Meet Ahuja
- Parul

ACADEMIC CONTEXT

Course: Data and Visual Analytics (DVA) Institute: Newton School of Technology

Faculty Mentor: Satyaki Das

Submission Date: 18 - February 2024

Table Of Contents

1. Executive Summary
2. Sector & Business Context
3. Problem Statement & Objectives
4. Data Description
5. Data Cleaning & Preparation
6. KPI & Metric Framework
7. Exploratory Data Analysis (EDA)
8. Advanced Analysis
9. Dashboard Design
10. Insights Summary
11. Recommendations
12. Impact Estimation
13. Limitations
14. Future Scope
15. Conclusion
16. Appendix
17. Contribution Matrix

1. Executive Summary:

The Problem: Beyond the "Accident"

Road safety is often treated as a series of random, unfortunate events. However, our analysis of traffic data reveals that accidents are rarely random; they are the result of specific, recurring risk factors.

With rising urban density, the goal of this project is to move from reactive policing to proactive intervention.

We address three main questions: where are the geographic hotspots, which driver demographics are most at risk, and how do environmental conditions (like road surface and lighting) turn a minor error into a fatal collision?

This project decodes the complexity of traffic incidents by addressing three fundamental pillars: the identification of geographic hotspots where infrastructure fails to support volume, the isolation of high-risk driver demographics, and the analysis of how environmental "severity multipliers"—such as poor lighting and road surface conditions—transform minor human errors into fatal collisions. By quantifying these variables, the report provides a data-driven roadmap to move beyond treating symptoms and begin addressing the root causes of road trauma.

Executive-Brief

Highlights

=> This project shifts road safety from a reactive "accident report" model to a proactive "risk management" model by analyzing 32 variables across human, vehicle, and environmental factors.

=> Using Google Sheets analytical engine, we identified that accidents are not random; they correlate heavily with specific `types_of_junction`, `light_conditions`, and driver `driving_experience` levels.

=> Our interactive dashboard allows stakeholders to use real-time filters (Slicers) to visualize "High-Risk Time Windows" (`time_category`), enabling better deployment of traffic personnel.

=> The analysis reveals that the severity of an accident is often more dependent on infrastructure (`lanes_or_medians`) and environmental triggers than on the `type_of_vehicle` alone.

=> We propose targeted infrastructure maintenance and experience-based safety campaigns, which are estimated to improve service efficiency and significantly reduce the `number_of_casualties`

Our Approach:

A Data-Driven Map to Safety

To solve this, we utilized a comprehensive dataset of traffic incidents, processing it entirely within Google Sheets.

Our methodology followed a four-step funnel:

1. Data Refinement: We cleaned and "engineered" features from messy raw data, such as categorizing `service_year_of_vehicle` and `driving_experience` to find correlations with `accident_severity`.

Metric Framework: We established KPIs like **Casualty-to-Vehicle Ratio** and **High-Risk Time Windows** to measure the true "danger level" of different scenarios.

2. Visual Analytics: We built a dynamic dashboard using **Pivot Tables and Slicers** to allow stakeholders to filter by weather, vehicle type, and area.

3. Root Cause Analysis: We looked past the surface to see if "Vehicle Defects" or "Road Alignment" were the silent killers behind the data.

The Experience Paradox: High-risk incidents aren't always caused by the youngest drivers; a significant spike occurs in specific "experience bands" where overconfidence may lead to riskier `vehicle_movement`.

- Environmental Triggers: While most accidents occur in "Good" weather, the severity of accidents increases drastically under specific `light_conditions` and `road_surface_type` combinations.
- Infrastructure Hotspots: Certain `types_of_junction` and `lanes_or_medians` configurations act as "bottlenecks of risk," showing a disproportionate number of `type_of_collision` incidents.
- Vehicle Condition: We identified a clear link between `service_year_of_vehicle` and the likelihood of mechanical failure leading to an accident.

2. Sector & Business Context

Sector Overview

The transport and logistics sector acts as the backbone of the modern economy, facilitating the movement of people and goods across urban and rural landscapes.

Within this sector, Road Safety Management has emerged as a critical sub-domain for both government transport departments and private insurance providers.

This sector relies heavily on "Telematics" and historical incident data to understand how human behavior, vehicle health, and infrastructure design interact on the road.

Current Challenges

Despite advancements in vehicle technology, the sector faces several persistent hurdles:

- Data Fragmentation: Information regarding `weather_conditions`, `road_surface_type`, and `vehicle_movement` is often recorded in silos, making it difficult to see the "big picture" of an accident.
- Reactive vs. Proactive Measures: Most interventions happen after an area is already labeled a "blackspot," rather than using predictive analytics to identify risks before they lead to fatalities.
- Human Element Complexity: Understanding the correlation between `driving_experience` and `casualty_severity` remains a challenge, as driver behavior is often unpredictable and influenced by external factors like `light_conditions`.

Aging Infrastructure: Many road layouts, such as specific `types_of_junction` or `lanes_or_medians`, were not designed for the current volume of heavy `type_of_vehicle` traffic

Why This Problem Was Chosen

This specific problem was selected for the DVA Capstone because it offers a perfect blend of social impact and analytical complexity:

- Rich Multidimensional Data: The dataset provides "next-level" info including technical details (`Defect_of_vehicle`), environmental data (`weather_conditions`), and demographic data (`age_band_of_driver`), allowing for deep exploratory analysis.
- Measurable Success Criteria: Unlike abstract business problems, road safety has clear KPIs—such as reducing `number_of_casualties`—making the impact of our data-driven recommendations easy to quantify.
- Practical Application of Google Sheets: This project demonstrates that high-level analytical insights (like `Correlation analysis` and `Trend analysis`) can be achieved using accessible tools like Google Sheets without needing expensive specialized software.
- Policy Relevance: The insights derived from `area_accident_occurred` and `cause_of_accident` can be directly translated into policy recommendations for local transport authorities.

4. Problem Statement & Objectives

Formal Problem Definition

The core challenge addressed in this report is the high frequency and severity of road traffic accidents. While accidents are often viewed as isolated incidents, they are driven by systemic risks involving driver demographics (e.g., `age_band_of_driver`), environmental factors (e.g., `weather_conditions`), and mechanical integrity (e.g., `Defect_of_vehicle`). This project aims to analyze these multidimensional data points to identify "High-Risk Profiles" and "Accident Hotspots". By doing so, we shift from a reactive stance to a proactive safety model that can predict where and why the next major incident is likely to occur.

Project Scope

To ensure the analysis remains focused and actionable, the project scope is defined as follows:

- Data Source: Use of the specific traffic accident dataset containing 32 distinct columns, ranging from `time_hour` to `cause_of_accident`.
- Tools: All data cleaning, primary transformation, and dashboarding are strictly executed within Google Sheets.
- Analysis Depth: The study focuses on Trend, Comparison, Distribution, and Correlation analysis.
- Inclusions: Analyzing the relationship between vehicle types, road surface conditions, and casualty severity.
- Exclusions: This project does not include real-time GPS tracking or post-accident insurance claim processing.

Success Criteria

The project will be considered successful if the following benchmarks are met:

- Identification of Key Risk Drivers: Clearly identifying the top 3 factors (e.g., specific `age_band_of_driver` or `road_alignment`) that contribute most to `fatal` or `serious` accident severity.
- Actionable Dashboard: Delivery of a functional Google Sheets dashboard with interactive filters (Slicers) that allows a non-technical user to identify hotspots by `area_accident_occurred`.

Evidence-Based Recommendations: Providing at least 5 policy or business recommendations that are directly mapped to data insights and have a clear "Business Impact" estimation.

Data Integrity: Achieving a 100% clean dataset where `missing values` and `outliers` are handled according to the documented methodology.

5. Data Description

5.1 Exact Dataset Source Citation & Access Link

- Source: Road Traffic Accident (RTA) Dataset.
- Citation: The dataset is a benchmark collection of road traffic accidents typically sourced from police records or road safety research repositories
- Access Link: RTA Dataset - Kaggle (<https://www.kaggle.com/datasets/saurabhshahane/road-traffic-accidents/data>)

5.2 Data Structure

The dataset follows a flat-file schema where each row represents a unique accident event, and columns represent the multidimensional attributes of that event.

Metadata Attribute	Specification
Data Architecture	Wide Format (Single-row-per-incident)
Total Record Count	12,316 Observations
Raw Dimensionality	32 Features
Post-Cleaning Dimensionality	34 Features (including engineered categories)
Primary Key	Time + Day_of_week (Composite logic for identification)
Storage Format	.csv (Comma Separated Values)
File Environment	Google Sheets Enterprise

5.3 In-Depth Column Explanation (Data Dictionary)

Below is the comprehensive mapping of the columns used in the analysis. This table defines the nature of the data you are processing

Column Name	Description
age_band_of_driver	Age bracket of the vehicle operator at the time of incident.
sex_of_driver	Gender identification of the driver.
educational_level	Highest formal education level achieved by the driver.
driving_experience	Total years of licensed driving history.
driving_relation	Relationship of the driver to the vehicle (Owner, Employee, etc.).
type_of_driver	Professional classification (e.g., Regular, Employee, Professional).

5.6 Data Limitations

1. Imbalance: The heavy skew toward "Slight" accidents makes it difficult to find statistical significance for "Fatal" accidents without advanced filtering.
2. Qualitative Bias: Fields like [Cause_of_accident](#) are subjective and depend on the reporting officer's interpretation.
3. No Location Coordinates: The dataset provides "Area Type" but lacks Latitude/Longitude, preventing high-resolution GIS Heatmap plotting.
4. Temporal Gaps: There is no "Date" column (only Day of Week), preventing Year-over-Year (YoY) trend analysis.

6. Data Cleaning & Preparation

The objective of this phase was to convert raw, observational incident data into a "Golden Record" suitable for high-fidelity visual analytics. All operations were executed within Google Sheets.

6.1 Structural & Schema Refinement

To ensure programmatic compatibility for complex formulas, the dataset underwent a structural overhaul.

Header Standardization: All original headers were converted to `snake_case`

(e.g., Type of vehicle became `type_of_vehicle`) to ensure ease of use in Google Sheets named ranges.

- Dimensional Reduction: Five columns with high cardinality or low analytical relevance were removed to optimize Google Sheets performance: `owner_of_vehicle`, `vehicle_driver_relation`, `work_of_casualty`, `fitness_of_casualty`, and `sex_of_casualty`.

6.2 Data Cleansing & Normalization

A rigorous cleansing pass was performed to handle noise, typos, and formatting inconsistencies.

Task	Operation	Impact / Result
Deduplication	Record-level scan	1 duplicate record was identified and removed to ensure data integrity.
Sanitization	Whitespace Trimming	Applied <code>=TRIM()</code> to all categorical columns to prevent "Ghost Categories" during pivot table aggregation.
Typo Correction	Fuzzy Matching	Corrected systematic errors such as "statioUnknownry" to "stationary" (481 rows).
Value Mapping	Formatting	Standardized vehicle service years by replacing variations like "yrs" and "yrss" with "years".
Placeholder Fix	Character Replacement	Replaced "?" placeholders with "Unknown" in the <code>type_of_vehicle</code> column to maintain categorical consistency.

6.3 Missing Values Handling

The dataset contained significant gaps, particularly in technical and demographic fields. These were addressed through logical imputation.

- Imputation with "Unknown": For critical columns where the value could not be inferred, blanks were replaced with "Unknown" to maintain row integrity. This included 1,548 rows in `driver_age_band`.
- Assumption-Based Logic: For `defect_of_vehicle`, 4,429 blank values were identified. Based on the assumption that a lack of reported defect implies a functional vehicle, these were imputed as "No defect".
- Critical Data Gaps: Identified a high systemic entry failure (4,443 rows) in `casualty_class`, `sex_of_casualty`, and the target variable `accident_severity`, which were flagged for careful treatment

during analysis.

6.4 Feature Engineering & Encoding

New attributes were engineered using Google Sheets formulas to facilitate temporal and numerical analysis.

- Temporal Binning: To analyze cyclical patterns, the raw `time` string was binned into four shifts:
 - Night: 00:00 - 05:59
 - Morning: 06:00 - 11:59
 - Afternoon: 12:00 - 17:59
 - Evening: 18:00 - 23:59
- Categorical Grouping: Reduced noise in `educational_level` by aggregating tiers into Low Education, School Education, and Higher Education.

Boolean Risk Flag: Created `vehicle_condition_risk` to identify "Risky" vs "Normal" vehicles based on the `Defect_of_vehicle` column.

6.5 Technical Formula Registry

The following core formulas were implemented to maintain a dynamic cleaning pipeline:

- Temporal Binning:
`=IF(HOUR(A2)<6,"Night",IF(HOUR(A2)<12,"Morning",IF(HOUR(A2)<18,"Afternoon","Evening")))`
- Gender Encoding (Array):
`=ArrayFormula(IF(D2:D="Male", 1, D2:D="Female", 0, D2:D="Unknown","Unknown"))`
- Vehicle Risk Logic:
`=ARRAYFORMULA(IF(J2:J="No defect", "Normal", IF(J2:J="Unknown", "Unknown", "Risky")))`

6. KPI & Metric Framework

The following metrics were engineered to transform raw incident counts into actionable safety intelligence.

6.1 Volume & Baseline Metrics

KPI 1: Cumulative Casualty Volume

- Value: 19,067
- Technical Definition: The absolute total count of individuals who sustained an injury or fatality across the 12,316 recorded incidents.
- Formula: =SUM(number_of_casualties)
- Analytical Significance: Establishes the gross human impact. It serves as the "Primary Anchor" for the study, providing a scale against which all subsequent risk factors are measured.

6.2 Severity & Distribution Metrics

KPI 2: Injury Severity Profile

- Value: 84.56% Slight | 14.15% Serious | 1.28% Fatal
- Technical Definition: The proportional distribution of all accidents categorized by the highest degree of physical harm sustained.
- Methodology: Pivot Table aggregation using % of Column Total.
- Analytical Significance: Shifts the focus from "if" an accident happens to "how lethal" it is. It isolates the specific variables contributing to the 1.28% fatal threshold, which represents the highest-value target for intervention.

6.3 Dynamic & Interaction Metrics

KPI 3: Average Vehicle Entanglement

- Value: 2.04
- Technical Definition: The mean number of vehicles participating in a single collision event.
- Formula: =AVERAGE(number_of_vehicles_involved)
- Analytical Significance: Defines the crash dynamic. A value \$>2.0\$ confirms that accidents are predominantly multi-vehicle "interaction errors" (such as junction failures or lane violations) rather than isolated single-vehicle mechanical failures.

KPI 4: Impact Intensity Ratio (Occupant Vulnerability)

- Value: 75.86%
- Technical Definition: The average number of casualties produced per vehicle involved in a crash.
- Formula: =SUM(number_of_casualties) / SUM(number_of_vehicles_involved)
- Analytical Significance: Measures the "Efficiency of Harm." It evaluates occupant vulnerability and helps identify if specific vehicle types (e.g., public transport) carry a disproportionately high human toll per incident.

6.4 Advanced Risk Metrics

KPI 5: Peak Severity Index (PSI)

- Value: 2.16
- Technical Definition: The maximum calculated ratio of casualties to crashes identified within a specific environmental or temporal segment (e.g., Night-time).
- Formula: MAX(SUM(number_of_casualties) / COUNTA(accident_severity)) grouped by a specific dimension.
- Analytical Significance: Pinpoints the "Crisis Point." It reveals that under specific worst-case scenarios, a single crash results in 2.16 casualties—nearly double the dataset average—providing a clear mandate for targeted policy enforcement.

8. Exploratory Data Analysis (EDA)

Using the KPIs defined above, this section analyzes the patterns discovered during the data exploration phase in Google Sheets.

8.1 Distribution Analysis: The Hierarchy of Harm

The primary investigation focused on the target variable: `accident_severity`.

Finding: 84.56% of incidents resulting in Slight injuries, 14.15% in Serious injuries, and only 1.28% being Fatal.

Insight: Safety interventions must look beyond high-frequency "Slight" accidents to identify the rare but catastrophic conditions that lead to the 1.28% fatal rate.

8.2 Comparison Analysis: Human Impact vs. Vehicle Involvement

With an Average of 2.04 vehicles per accident, the data suggests that traffic safety is largely an issue of interaction rather than single-driver error hitting stationary objects.

The 75.86% Rule: Every vehicle involved in a crash has a nearly 76% probability of producing at least one casualty.

- High-Stakes Segments: Preliminary analysis shows that the Casualty-to-Vehicle Ratio increases significantly in "Public Transport" categories compared to private automobiles, likely due to higher occupancy rates.

8.3 Peak Risk Discovery

The Highest Severity Index (2.16) was identified during specific "Night" and "Evening" shifts.

Finding: Under peak-risk conditions, a single crash results in more than double the average casualty rate.

Strategic Implication: This confirms that nighttime visibility (Light Conditions) and driver fatigue (Temporal Patterns) are the most significant multipliers of accident

9. Advanced Analysis

Moving beyond basic counts to identify systemic risk multipliers and behavior patterns.

- Segmentation (Temporal Risk Profiling): Analyzed the relationship between traffic volume and incident lethality. While accidents are most frequent during the day, segmentation revealed a "Severity Surge" at night, with the Severity Index peaking at 2.16. This identifies darkness as a primary force multiplier for fatal outcomes.
- Root Cause Analysis (The Experience Paradox): Conducted a comparative study of driver experience vs. injury level. Found that while junior drivers (<2yr) have higher accident frequency, experienced drivers (>10yr) are disproportionately involved in higher-severity (Fatal/Serious) collisions, suggesting complacency-linked risk.
- Risk/Anomaly Detection: Identified "Office Areas" as geographic anomalies where accidents do not follow the standard "Weekend Peak" trend, showing a significant surge during weekday rush hours (8 AM – 10 AM).

10. Dashboard Design

The project features an interactive High-Visibility command center built in Google Sheets to facilitate data-driven decision-making.

- Dashboard Objective: To provide transport authorities with a real-time tool to isolate "High-Lethality" variables and visualize the scale of human impact (19,067 casualties).
- View Structure:
 - Executive Scorecard: Top-level row featuring the 5 primary KPIs (Casualties, Fatality %, Avg Vehicles, Ratio, and Max Severity).
 - Strategic Layer: Middle section containing the Junction Type distribution and Environmental condition analysis.
 - Operational Layer: Bottom section focusing on Driver Demographics and Vehicle Metadata.

- Filters & Drilldowns:
 - * Global Slicers: Interactive filters for Day_of_Week, Road_Surface, and Light_Conditions.
 - Dynamic Range: Formulas automatically update the Severity Index based on the filtered segment.

11. Insights Summary

Strategic findings derived from the multi-dimensional analysis of the RTA Dataset.

Demographic High-Risk Profile: Male drivers in the "Young" to "Middle-aged" bands are the primary contributors to casualty volumes, representing a critical demographic for safety intervention.

The Experience Paradox: While "Junior" drivers (<2 years) cause frequent slight injuries due to lack of skill, "Expert" drivers (>10 years) show a plateau in accident involvement, suggesting that overconfidence/complacency is as dangerous as inexperience.

Infrastructural Bottlenecks: "Office Areas" have emerged as the highest-density hotspots, indicating that peak-hour transit designs in corporate hubs are currently insufficient for the actual traffic volume.

Behavioral Dominance: Over 75% of accidents are rooted in "Lane Violations" and "Reckless Driving," proving that human intervention and enforcement are more urgent than mechanical inspections.

The Counter-Intuitive Visibility Risk: Most accidents occur during "Daylight" and "Dry" conditions. Optimal environmental factors lead to higher speeds and reduced driver caution, creating a paradox where "safe" conditions result in more frequent collisions.

Multi-Vehicle Dynamics: With an average of 2.04 vehicles per accident, the data confirms that "Side-swipe" and "Rear-end" collisions are the dominant crash types, pointing to poor following-distance discipline.

Night-Time Severity Surge: Although night-time volume is lower, the Severity Index (2.16) is significantly higher, indicating that crashes occurring in darkness are twice as likely to result in multiple casualties.

Vulnerability of Public Transport: Segmentation shows that while lorries and buses are fewer in number, their Casualty-to-Vehicle Ratio is disproportionately high, making them "High-Stakes" vehicles in the safety ecosystem.

12. Recommendations

Actionable strategies mapped to the insights discovered during analysis.

=> Pillar 1: Infrastructure—The "Office Zone" Initiative

- Target Area: High-density Office and Commercial zones.
- Data Evidence: These areas account for the highest concentration of weekday incidents.
- Action: Implement "Smart Speed Calming" (Speed humps/narrowing) and dedicated pedestrian "Refuge Islands" in high-frequency office corridors.
- Expected Impact: High | Reduces pedestrian-to-vehicle collision severity.

Pillar 2: Behavior—The "Lane Discipline" AI Enforcement

- Target Cause: Lane violations and improper overtaking (75% Human Error).
- Data Evidence: Driver behavior is the #1 behavioral trigger for casualties.
- Action: Deploy AI-integrated traffic cameras to automatically detect and fine lane-crossing violations, specifically during peak hours.
- Expected Impact: Very High |

Addresses the primary root cause of multi-vehicle collisions.

Pillar 3: Environmental—The "After-Hours" Visibility Audit

- Target Condition: Nighttime and Poor Lighting.
- Data Evidence: The Severity Index (2.16) peaks during darkness, making night accidents twice as lethal.
- Action: Conduct a lighting audit of "A-Class" roads and transition to high-luminosity LED street lighting to eliminate blind spots.
- Expected Impact: High | Directly lowers the fatality rate associated with low visibility.

Pillar 4: Mechanical—The "Vehicle Compliance" Mandate

- Target Risk: Mechanical defects and aging vehicle fleets.
- Data Evidence: Higher casualty-to-vehicle ratios (75.86%) are seen in older or commercial vehicle types.
- Action: Introduce mandatory quarterly safety audits for commercial lorries and public transport vehicles operating in high-risk zones.
- Expected Impact: Medium | Reduces risk from mechanical failures like brake or light defects.

13. Impact Estimation

This section quantifies the potential real-world value of implementing the recommendations derived from our analysis.

- Save Cost (Economic Impact): Traffic accidents impose a massive financial burden on healthcare and insurance systems. By targeting the "Human Error" root causes (75% of our data), a projected **10% reduction in total accidents** could save an estimated **\$1.2M - \$2.5M annually** in emergency response costs and public property repairs.
- Improve Efficiency (Resource Allocation): Currently, patrols are distributed generally. By using our **Temporal Risk Windows**, law enforcement can shift 30% more personnel to the "Afternoon" and "Evening" shifts where volume peaks, improving response times without increasing headcount.
- Improve Service (Public Safety): Addressing infrastructure in "Office Area" hotspots (our primary geographic finding) will reduce commute delays by an estimated 15%, improving the overall quality of life for urban workers.

14. Limitations

Every data project has boundaries; identifying them ensures the credibility of our conclusions.

- Data Issues (Imbalance): The dataset is heavily skewed toward "Slight" accidents (84.56%). This small sample size of "Fatal" incidents (1.28%) makes it statistically difficult to build a definitive predictive model for deaths compared to minor injuries.
- Assumption Risks: Our analysis assumes that "Unknown" values in the `defect_of_vehicle` column represent a functional vehicle. If these "Unknowns" actually hide mechanical failures, our "Human Error" dominance might be slightly over-represented.
- What Cannot Be Concluded: Due to the lack of **Latitude/Longitude (GPS)** coordinates, we cannot pinpoint specific "Black Spots" at the street level; we can only identify general zones like "Office" or "Residential" areas.

15. Future Scope

How this project can evolve into a more powerful decision-making tool.

- Advanced Predictive Modeling (Python/R): Moving beyond Google Sheets to implement a Random Forest or XGBoost model to predict accident severity in real-time based on live weather feeds.
- Integration of External Data: Incorporating Telematics Data (GPS tracking from public buses) and Traffic Flow Sensors to correlate accident density with actual road congestion levels.
- GIS Mapping: Using Tableau or Power BI to create interactive Heatmaps that allow city planners to zoom into specific intersections.
- NLP on Police Reports: Using Natural Language Processing to analyze the text descriptions of accidents to find hidden patterns not captured in categorical columns.

16. Conclusion

Final summary of value delivered.

This project successfully transformed a raw dataset of 12,316 records into a strategic road safety roadmap. By executing a rigorous cleaning pipeline and advanced segmentation in Google Sheets, we identified that road safety is not just a matter of "bad luck" but a predictable intersection of demographic behavior (Young Males), temporal shifts (Evening peaks), and environmental overconfidence (Daylight/Dry conditions).

The insights provided—ranging from the Experience Paradox to Office-Zone hotspots—deliver a data-driven foundation for policy makers to move away from reactive policing and toward proactive, evidence-based safety

interventions. The ultimate value delivered is a framework that prioritizes human life through analytical precision.

17. Appendix

17.1 Extended Data Dictionary

[This section provides the full technical schema for the 33 variables processed from the raw source data.](#)

- Primary Key: [Accident_Index](#) (Unique alphanumeric identifier for each incident).
- Target Variable: [Accident_Severity](#) (Ordinal: 1=Slight, 2=Serious, 3=Fatal).
- Temporal Features: [Day_of_Week](#), [Time](#), [Time_Category](#) (Derived).

Environmental Features: [Road_Surface](#), [Light_Conditions](#),

- Weather_Conditions, Special_Conditions_at_Site.
- Human Factors: Age_band_of_driver, Sex_of_driver, Driving_experience, Educational_level.
- Vehicle Metadata: Type_of_vehicle, Owner_of_vehicle, Service_year_of_vehicle, Defect_of_vehicle.

18. Contribution Matrix (Mandatory)

Team Member	Dataset & Sourcing	Cleaning	KPI & Analysis	Dashboard	Report Writing	PPT	Overall Role
Ganga	yes	yes	No	yes	yes	yes	
Krishnadevann	yes	yes	yes	yes			
Siddhant	No	Yes		Yes	No	No	
Arjun	No	yes	yes	No	No	yes	
Meet Ahuja	No	No	Yes	Yes	Yes	No	
Parul	No	No	No	No	No	No	